

Table of contents

2025-02-03 . . . . .	2
2025-03-04 . . . . .	2
2025-03-19 . . . . .	2

**2025-02-03**

- Walked Mark through the in-progress prototype for the claims-based frame extraction (Document will be uploaded later this week).
- Current state of the progress:
  - Take the ABC demo dataset.
  - Use LLM to extract claims made about energy sources.
  - Use a sentiment model to detect the ‘vibe’ of the claim.
  - Compare these claims to the distribution of the overall discourse to determine if the article frames a piece in a way that is consistent with the ‘ground truth’.
  - If it doesn’t, then we can also see if it is being contrarian (both polarised and biased against renewables) or presenting a false neutrality (when presenting both sides when one side is not truly present in the ground truth).
- Improvements that I want to make before next fortnight:
  - Multidimensional sentiment analysis: Checking sentiment across multiple concepts, such as cost, harm, effectiveness.
  - Add in additional energy infrastructure.
  - Add in a single score based on the probability the article follows the distribution in the ground truth.
- Will send a better explainer around this week, would love feedback from Simon and Libby

**2025-03-04**

- I have the demo for the ‘claims based extractor’ done, now I just need to get the dataset working.
- From the feedback from Simon, I agree that this is a bit reductionist, but I think it is less reductionist than most of the other work.
  - In response to this, I do want to expand the number of dimensions that I am looking at evaluations for. Currently, I do kind of just have an economic and ‘other’ dimension, but I think I can do better.
  - This is something that I want to get out after I set the dataset, because I think these will be better as emergent properties rather than things that I think of myself.
- Some other concerns about the current method:
  - Is the framing independent? I know that the probability that one claim occurs will not be independent of the other claims, but does this matter and does it impact it enough to lose the nice explainability?
  - Are there things not being picked up by the extractor?
- For next time, lets get the dataset (meeting with librarian to discuss).
- Present findings to the larger supervision group. Put together a bit of an intermediary commencement document

**2025-03-19**

Full meeting with the whole gang

- Started with a quick summary of the demo work, simon raised a couple of concerns re ‘The Rig’:
  - Shouldn’t be asking the language model to generate numbers, this should be the task of a logistical model ( Thinking Vader for training Data )

- Filter object could be as simple as a Black list white list for bigrams -> don't need to do llms.
- Discussed the concerns raised by the copyright expert:
  - We know that this is going to be an issue, and for some journals, they won't even publish if the data is scraped.
  - Factiva ect is not an option for large scale collection. The terms and condition say no to this.
  - Trove could be an option
  - The cost of getting a licence or an agreement with the news organisations will cost way too much for the scope of the PhD.
  - This has the potential for a good paper -> What are the current legal barriers for text mining in Australia.
- As for goals moving forward:
  1. Investigate / Write Up the legal barriers
  2. Write Up A New Proposal, focus on the why we are doing this:
    - a. Use as an intervention down the line
    - b. Investigate the overton-widow shifting events
  3. Create a test set for each of the portions of the rig and have a statistic for determining the efficacy of the process.